

## Data Ingestion from the RDS to HDFS using Sqoop

### Sqoop Import command used for importing table from RDS to HDFS:

Before running the Sqoop import command, I made sure the target directory is not already created. Otherwise, the Sqoop import command would throw an error:

```
hadoop fs -rm -r /user/root/SRC_ATM_TRANS
```

### Sqoop Import Command:

```
sqoop import \  
--connect jdbc:mysql://upgraddetest.cyaie1c9bmnf.us-east-1.rds.amazonaws.com/testdatabase \  
--table SRC_ATM_TRANS \  
--username student --password STUDENT123 \  
--target-dir /user/root/SRC_ATM_TRANS \  
-m 1
```

In the screenshot below, I can see that as a result of Sqoop Import Job, 2468572 records have been retrieved (same as the checkpoint mentioned in the Validation document):

```

21/08/09 18:16:28 INFO mapreduce.Job: map 0% reduce 0%
21/08/09 18:17:04 INFO mapreduce.Job: map 100% reduce 0%
21/08/09 18:17:13 INFO mapreduce.Job: Job job_1628531322983_0001 completed successfully
21/08/09 18:17:14 INFO mapreduce.Job: Counters: 30
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=176686
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=87
        HDFS: Number of bytes written=531214815
        HDFS: Number of read operations=4
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Other local map tasks=1
        Total time spent by all maps in occupied slots (ms)=41548
        Total time spent by all reduces in occupied slots (ms)=0
        Total time spent by all map tasks (ms)=41548
        Total vcore-milliseconds taken by all map tasks=41548
        Total megabyte-milliseconds taken by all map tasks=42545152
    Map-Reduce Framework
        Map input records=2468572
        Map output records=2468572
        Input split bytes=87
        Spilled Records=0
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=249
        CPU time spent (ms)=28410
        Physical memory (bytes) snapshot=419827712
        Virtual memory (bytes) snapshot=2802429952
        Total committed heap usage (bytes)=384827392
    File Input Format Counters
        Bytes Read=0
    File Output Format Counters
        Bytes Written=531214815
21/08/09 18:17:14 INFO mapreduce.ImportJobBase: Transferred 506.6059 MB in 64.3895 seconds (7.8678 MB/sec)
21/08/09 18:17:14 INFO mapreduce.ImportJobBase: Retrieved 2468572 records.
[root@ip-10-0-0-166 ~]#

```

### Command used to see the list of imported data in HDFS:

```
hadoop fs -ls /user/root/SRC_ATM_TRANS
```

In the screenshot below, I can see that the target directory contains 2 items:

- The first file is the success file, indicating that the MapReduce job was successful.
- The second file 'part-m-00000' is the one with all of the data I imported. Since I used only one mapper in my import command thus the data is in a single file.

```
[root@ip-10-0-0-166 ~]# hadoop fs -ls /user/root/SRC_ATM_TRANS
Found 2 items
-rw-r--r--  3 root supergroup          0 2021-08-09 18:17 /user/root/SRC_ATM_TRANS/_SUCCESS
-rw-r--r--  3 root supergroup 531214815 2021-08-09 18:16 /user/root/SRC_ATM_TRANS/part-m-00000
[root@ip-10-0-0-166 ~]#
```

When I open the 'part-m-00000' file using the following command, I can see all of the data that has been imported:

```
hadoop fs -cat /user/root/SRC_ATM_TRANS/part-m-00000
```

**Screenshot of a portion of the imported data:**

2017, January, 4, Wednesday, 10, Active, 104, NCR, Intern ÆøsterÆ, ÆøsterÆ, 12, 4886, Aalborg, 275.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Inactive, 3, NCR, Ikast, RÆdhusstrÆdet, 12, 7430, 56.1001, 69, 7, 350, 0.000, 0, 800, Clear, Sky is Clear

2017, January, 4, Wednesday, 10, Active, 33, NCR, Vadum, Ellehammersvej, 43, 9430, 57.118, 9.008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Inactive, 54, NCR, Durup, Torvet, 4, 7870, 56.745, 8.949, DKK, 1001, 69, 7, 350, 0.000, 32, 802, Clouds, scattered clouds

2017, January, 4, Wednesday, 10, Active, 38, NCR, Hasseris, Hasserisvej, 113, 9000, 57.044, 95.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 69, NCR, Taars, Bredgade, 91, 9830, 57.385, 10.116, D, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 31, NCR, Slagelse, Mariendals Alle, 29, 4200, 55.396.380, 994, 74, 10, 350, 0.000, 12, 801, Clouds, few clouds

2017, January, 4, Wednesday, 10, Active, 62, Diebold Nixdorf, Terndrup, Bymidten, 2, 9575, 5und, 275.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 41, Diebold Nixdorf, Skagen, Sct. Laurentiivej, 3613939, Skagen, 272.280, 997, 86, 9, 20, 0.000, 90, 600, Snow, light snow

2017, January, 4, Wednesday, 10, Active, 13, NCR, SÆby, Vestergade, 3, 9300, 57.334, 10.51avn, 276.229, 1009, 100, 13, 15, 0.000, 76, 803, Clouds, broken clouds

2017, January, 4, Wednesday, 10, Active, 15, NCR, Vestre, Kastetvej, 36, 9000, 57.053, 9.905, 9, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 48, Diebold Nixdorf, BrÆnderslev, Algade, 4, 97, Vadum, 275.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 104, NCR, Intern ÆøsterÆ, ÆøsterÆ, 12, 86, Aalborg, 275.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 4, NCR, Svogerslev, BrÆnsager, 1, 4000, 55.634, 1, 994, 74, 10, 350, 0.235, 36, 500, Rain, light rain

2017, January, 4, Wednesday, 10, Active, 40, Diebold Nixdorf, Frederikshavn, Danmarksgade 7, 2621927, Frederikshavn, 276.229, 1009, 100, 13, 15, 0.000, 76, 803, Clouds, broken clouds

2017, January, 4, Wednesday, 10, Active, 69, NCR, Taars, Bredgade, 91, 9830, 57.385, 10.116, D, 08, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 39, NCR, Svenstrup, GodthÆbsvej, 14, 9230, 56.97rup, 275.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 49, NCR, Bindslev, NÆrrebro, 18, 9881, 57.541, 109, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 62, Diebold Nixdorf, Terndrup, Bymidten, 2, 9575, 5und, 275.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is clear

2017, January, 4, Wednesday, 10, Active, 28, NCR, LÆgstÆr, Æøsterbrogade, 8, 9670, 5y, 56.962, 9.258, 2617467, Logstor, 275.079, 1008, 69, 11, 3, 0.000, 8, 800, Clear, sky is cle